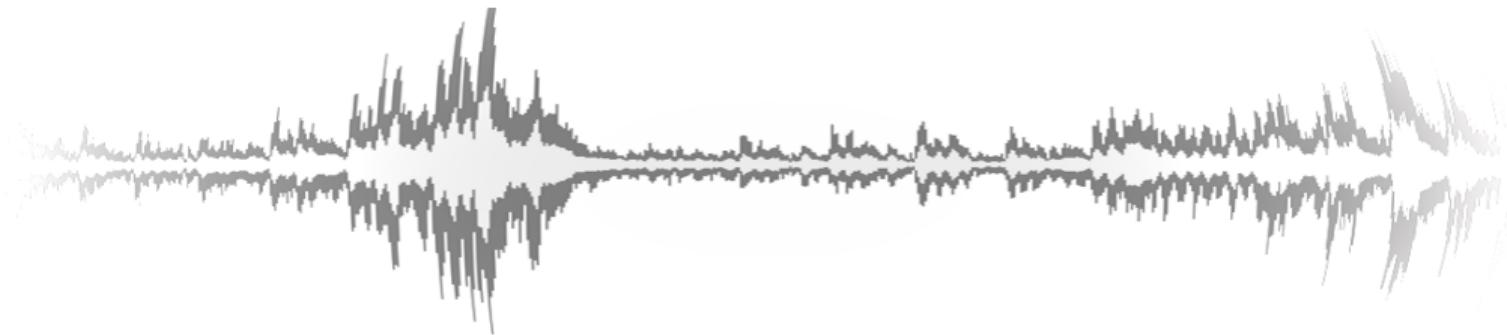


# Digital Signal Processing for Music

## Part 22: Time-stretching and Pitch-shifting

alexander lerch



# time stretching & pitch shifting

## introduction

- **time stretching**

change playback speed/tempo without changing pitch

- **pitch shifting**

change pitch without changing tempo/playback speed

- **terms**

- time/pitch scaling
- time expansion/compression

# time stretching & pitch shifting

## introduction

- **time stretching**

change playback speed/tempo without changing pitch

- **pitch shifting**

change pitch without changing tempo/playback speed

- **terms**

- time/pitch scaling
- time expansion/compression

# time stretching & pitch shifting

## introduction

- **time stretching**

change playback speed/tempo without changing pitch

- **pitch shifting**

change pitch without changing tempo/playback speed

- **terms**

- time/pitch scaling
- time expansion/compression

# time stretching & pitch shifting

## applications

- **beat matching:** align tempo of two or more audio files (mashup)
- **key lock:** “align” pitch of two or more audio files (mashup)
- **pitch/time correction:** edit intonation, frequency deviation, vibrato, glissando
- **video frame rate conversion**
- **sample player/libraries**
- **sound design**
- **educational software:** pitch and timing visualization

# time stretching & pitch shifting

## applications

- **beat matching:** align tempo of two or more audio files (mashup)
- **key lock:** “align” pitch of two or more audio files (mashup)
- **pitch/time correction:** edit intonation, frequency deviation, vibrato, glissando
- video frame rate conversion
- sample player/libraries
- sound design
- educational software: pitch and timing visualization

# time stretching & pitch shifting

## applications

- **beat matching:** align tempo of two or more audio files (mashup)
- **key lock:** “align” pitch of two or more audio files (mashup)
- **pitch/time correction:** edit intonation, frequency deviation, vibrato, glissando
- video frame rate conversion
- sample player/libraries
- sound design
- educational software: pitch and timing visualization

# time stretching & pitch shifting

## applications

- **beat matching:** align tempo of two or more audio files (mashup)
- **key lock:** “align” pitch of two or more audio files (mashup)
- **pitch/time correction:** edit intonation, frequency deviation, vibrato, glissando
- **video frame rate conversion**
- sample player/libraries
- sound design
- educational software: pitch and timing visualization

# time stretching & pitch shifting

## applications

- **beat matching:** align tempo of two or more audio files (mashup)
- **key lock:** “align” pitch of two or more audio files (mashup)
- **pitch/time correction:** edit intonation, frequency deviation, vibrato, glissando
- **video frame rate conversion**
- **sample player/libraries**
- **sound design**
- **educational software:** pitch and timing visualization

# time stretching & pitch shifting

## applications

- **beat matching:** align tempo of two or more audio files (mashup)
- **key lock:** “align” pitch of two or more audio files (mashup)
- **pitch/time correction:** edit intonation, frequency deviation, vibrato, glissando
- **video frame rate conversion**
- **sample player/libraries**
- **sound design**
- **educational software:** pitch and timing visualization

# time stretching & pitch shifting

## applications

- **beat matching:** align tempo of two or more audio files (mashup)
- **key lock:** “align” pitch of two or more audio files (mashup)
- **pitch/time correction:** edit intonation, frequency deviation, vibrato, glissando
- **video frame rate conversion**
- **sample player/libraries**
- **sound design**
- **educational software:** pitch and timing visualization

# time stretching & pitch shifting

## stretch and pitch factors

$$s = \frac{t_{output}}{t_{input}}$$
$$p = \frac{f_{output}}{f_{input}}$$

### examples:

- half speed:  $s = 2$
- half pitch:  $p = \frac{1}{2}$
- semitone up/down:

$$p_u = 2^{1/12} = 1.059 \quad p_d = 2^{-1/12} = 0.9439$$

- 100 BPM → 75 BPM:  $s = \frac{4}{3}$

# time stretching & pitch shifting

## stretch and pitch factors

$$s = \frac{t_{output}}{t_{input}}$$
$$p = \frac{f_{output}}{f_{input}}$$

### examples:

- *half speed*:  $s = 2$
- *half pitch*:  $p = \frac{1}{2}$
- *semitone up/down*:

$$p_u = 2^{1/12} = 1.059 \quad p_d = 2^{-1/12} = 0.9439$$

- *100 BPM → 75 BPM*:  $s = \frac{4}{3}$

# time stretching & pitch shifting

## stretch and pitch factors

$$s = \frac{t_{output}}{t_{input}}$$
$$p = \frac{f_{output}}{f_{input}}$$

### examples:

- half speed:  $s = 2$
- half pitch:  $p = \frac{1}{2}$
- semitone up/down:

$$p_u = 2^{1/12} = 1.059 \quad p_d = 2^{-1/12} = 0.9439$$

- 100 BPM → 75 BPM:  $s = \frac{4}{3}$

# time stretching & pitch shifting

## stretch and pitch factors

$$s = \frac{t_{output}}{t_{input}}$$
$$p = \frac{f_{output}}{f_{input}}$$

### examples:

- *half speed*:  $s = 2$
- *half pitch*:  $p = \frac{1}{2}$
- *semitone up/down*:

$$p_u = 2^{1/12} = 1.059 \quad p_d = 2^{-1/12} = 0.9439$$

- *100 BPM → 75 BPM*:  $s = \frac{4}{3}$

# time stretching & pitch shifting

## stretch and pitch factors

$$s = \frac{t_{output}}{t_{input}}$$
$$p = \frac{f_{output}}{f_{input}}$$

### examples:

- *half speed*:  $s = 2$
- *half pitch*:  $p = \frac{1}{2}$
- *semitone up/down*:

$$p_u = 2^{1/12} = 1.059 \quad p_d = 2^{-1/12} = 0.9439$$

- *100 BPM → 75 BPM*:  $s = \frac{4}{3}$

# time stretching & pitch shifting

## stretch and pitch factors

$$s = \frac{t_{output}}{t_{input}}$$
$$p = \frac{f_{output}}{f_{input}}$$

### examples:

- *half speed:*  $s = 2$
- *half pitch:*  $p = \frac{1}{2}$
- *semitone up/down:*

$$p_u = 2^{1/12} = 1.059 \quad p_d = 2^{-1/12} = 0.9439$$

- *100 BPM → 75 BPM:*  $s = \frac{4}{3}$

# time stretching & pitch shifting

## stretch and pitch factors

$$s = \frac{t_{output}}{t_{input}}$$
$$p = \frac{f_{output}}{f_{input}}$$

### examples:

- *half speed:*  $s = 2$
- *half pitch:*  $p = \frac{1}{2}$
- *semitone up/down:*

$$p_u = 2^{1/12} = 1.059 \quad p_d = 2^{-1/12} = 0.9439$$

- *100 BPM → 75 BPM:*  $s = \frac{4}{3}$

# time stretching & pitch shifting

## stretch and pitch factors

$$s = \frac{t_{output}}{t_{input}}$$
$$p = \frac{f_{output}}{f_{input}}$$

### examples:

- *half speed:*  $s = 2$
- *half pitch:*  $p = \frac{1}{2}$
- *semitone up/down:*

$$p_u = 2^{1/12} = 1.059 \quad p_d = 2^{-1/12} = 0.9439$$

- *100 BPM → 75 BPM:*  $s = \frac{4}{3}$

time stretching & pitch shifting  
stretch and pitch factors

$$s = \frac{t_{output}}{t_{input}}$$

$$p = \frac{f_{output}}{f_{input}}$$

## examples:

- half speed:  $s = 2$
  - half pitch:  $p = \frac{1}{2}$
  - semitone up/down:

$$p_u = 2^{1/12} = 1.059 \quad p_d = 2^{-1/12} = 0.9439$$

- $100 \text{ BPM} \rightarrow 75 \text{ BPM}: s = \frac{4}{3}$

# time stretching & pitch shifting

## resampling

- **traditional:** resampling

- change inter-sample 'distance' by interpolation
- keep playback sample rate constant
- audio example



⇒ tempo change results in pitch change (and vice versa)

$$s = \frac{1}{p}$$

# time stretching & pitch shifting

## resampling

- **traditional:** resampling

- change inter-sample 'distance' by interpolation
- keep playback sample rate constant
- audio example
  - original
  - resample

⇒ tempo change results in pitch change (and vice versa)

$$s = \frac{1}{p}$$

# time stretching & pitch shifting

## resampling

- **traditional:** resampling

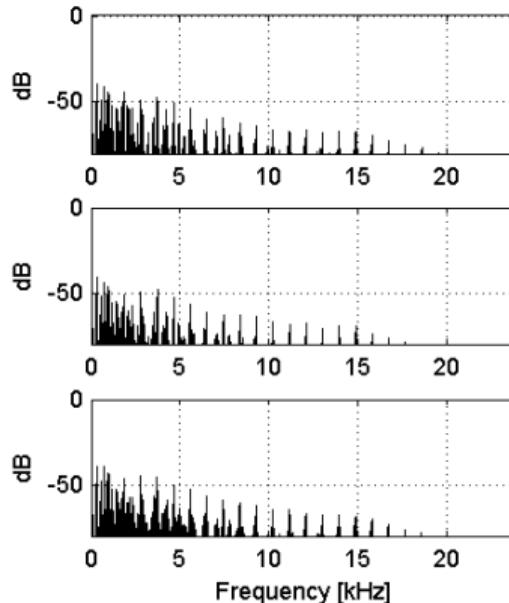
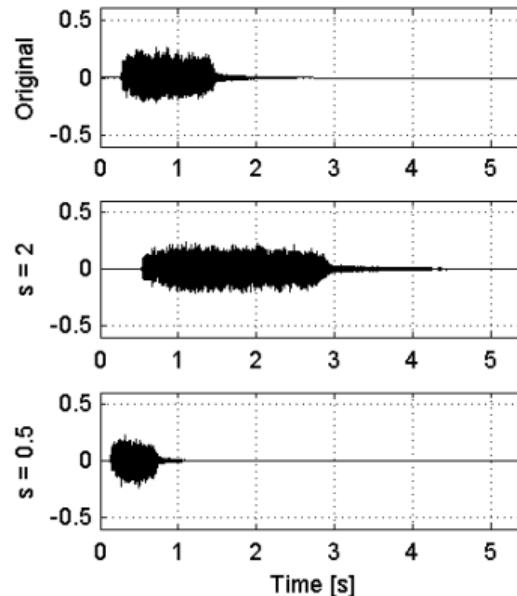
- change inter-sample 'distance' by interpolation
- keep playback sample rate constant
- audio example
  - original
  - resample

⇒ tempo change results in pitch change (and vice versa)

$$s = \frac{1}{p}$$

# time stretching & pitch shifting

## stretching: frequency domain



# OLA

## introduction

overlap and add approaches for, e.g.,

- granular synthesis
- time/frequency synthesis and processing
- time-stretching and pitch-shifting

# OLA

## introduction

overlap and add approaches for, e.g.,

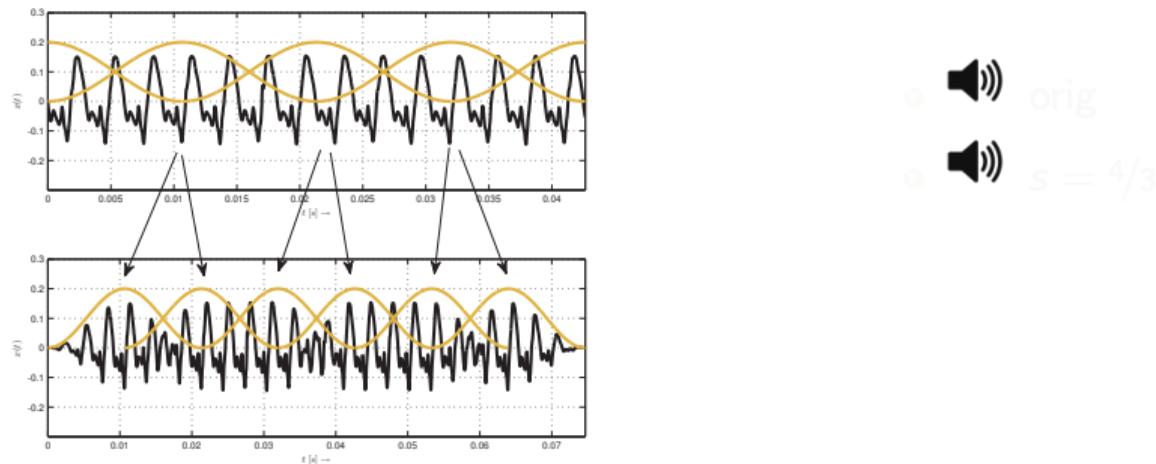
- granular synthesis
- time/frequency synthesis and processing
- time-stretching and pitch-shifting

# OLA

## time stretching

### overlap and add

- ➊ split input signal into overlapping blocks
- ➋ duplicate or discard blocks depending on stretch factor

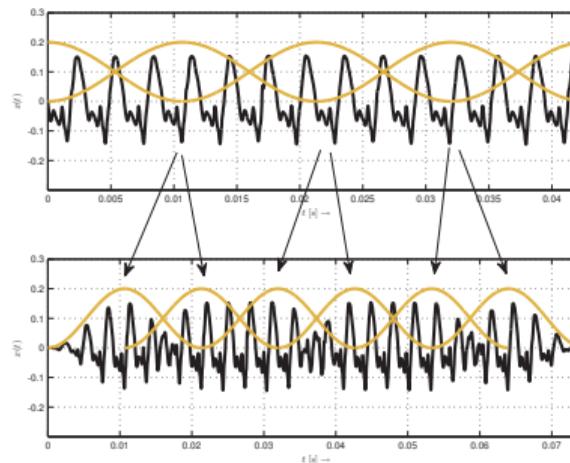


# OLA

## time stretching

### overlap and add

- ① **split input signal into overlapping blocks**
- ② **duplicate or discard blocks** depending on stretch factor



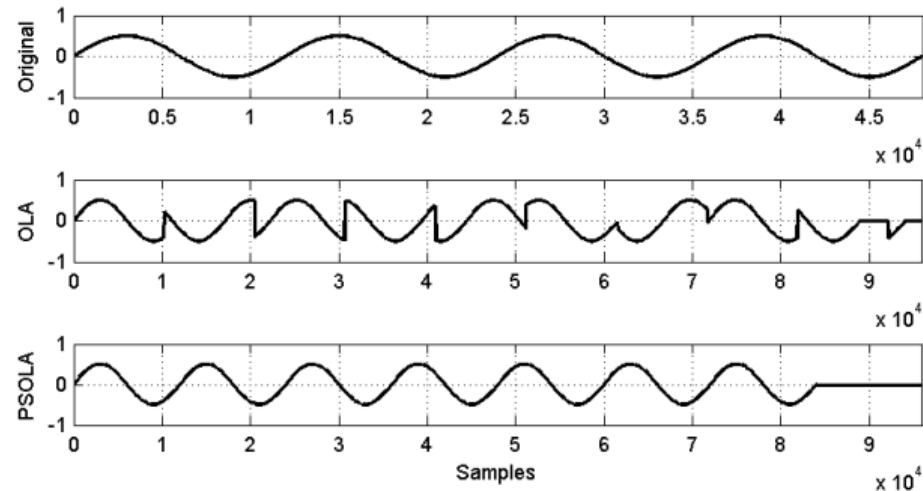
- 🔊 orig
- 🔊  $s = 4/3$

# PSOLA

## time stretching

### pitch synchronous overlap and add

- use the OLA principle, but
- adapt block length** to fundamental period length

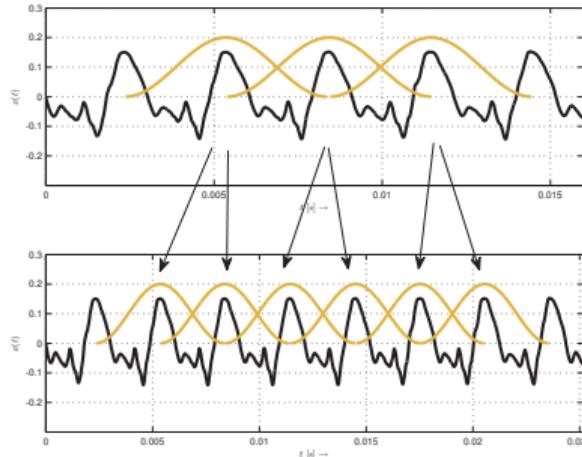


# PSOLA

## time stretching

### pitch synchronous overlap and add

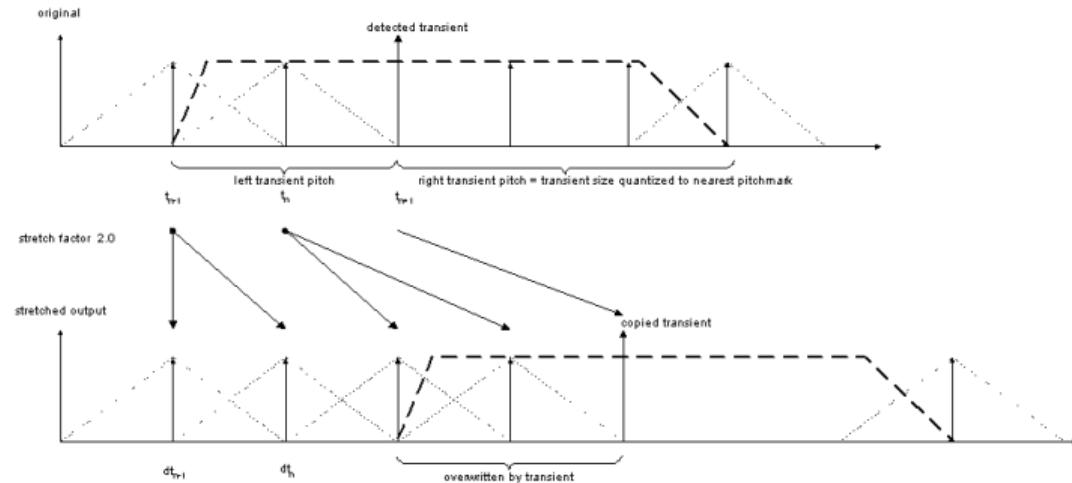
- use the OLA principle, but
- adapt block length** to fundamental period length



- orig
- $s = 4/3$

# PSOLA

## transient copying



# PSOLA

## time stretching summary

### ● processing steps

- ① detect *fundamental frequency/period length*
- ② set *pitch marks*
- ③ intelligently *select blocks* to be repeated/discard

### ● advantages

- *high granularity* — modify audio on period length resolution
- *high quality*

### ● problems

- quality depends on *pitch tracking* reliability
- quality and timbre depends on *pitch mark positioning*
- works only for *monophonic* input signals
  - polyphonic and noisy segments
  - reverberation and overlapping tones
- *noise, plosives* require special consideration
- *copying artifacts* (double transients, timing deviations)

### ● typical applications

- standard approach for vocal editing tools

# PSOLA

## time stretching summary

### ● processing steps

- 1 detect *fundamental frequency/period length*
- 2 set *pitch marks*
- 3 intelligently *select blocks* to be repeated/discard

### ● advantages

- *high granularity* — modify audio on period length resolution
- *high quality*

### ● problems

- quality depends on *pitch tracking* reliability
- quality and timbre depends on *pitch mark positioning*
- works only for *monophonic* input signals
  - polyphonic and noisy segments
  - reverberation and overlapping tones
- *noise, plosives* require special consideration
- *copying artifacts* (double transients, timing deviations)

### ● typical applications

- standard approach for vocal editing tools

# PSOLA

## time stretching summary

### processing steps

- 1 detect *fundamental frequency/period length*
- 2 set *pitch marks*
- 3 intelligently *select blocks* to be repeated/discard

### advantages

- *high granularity* — modify audio on period length resolution
- *high quality*

### problems

- quality depends on *pitch tracking* reliability
- quality and timbre depends on *pitch mark positioning*
- works only for *monophonic* input signals
  - polyphonic and noisy segments
  - reverberation and overlapping tones
- *noise, plosives* require special consideration
- *copying* artifacts (double transients, timing deviations)

### typical applications

- standard approach for vocal editing tools

# PSOLA

time stretching summary

## • processing steps

- 1 detect *fundamental frequency/period length*
- 2 set *pitch marks*
- 3 intelligently *select blocks* to be repeated/discard

## • advantages

- *high granularity* — modify audio on period length resolution
- *high quality*

## • problems

- quality depends on *pitch tracking* reliability
- quality and timbre depends on *pitch mark positioning*
- works only for *monophonic* input signals
  - polyphonic and noisy segments
  - reverberation and overlapping tones
- *noise, plosives* require special consideration
- *copying* artifacts (double transients, timing deviations)

## • typical applications

- **standard approach for vocal editing tools**

# time stretching

## phase vocoder

- ① **split input signal into overlapping blocks**
- ② **compute magnitude and phase spectrum of each block**
- ③ **change overlap ratio between blocks depending on stretch factor**
- ④ **keep the magnitude, adapt the phase per bin to the block's new time stamp**



# time stretching

## phase vocoder

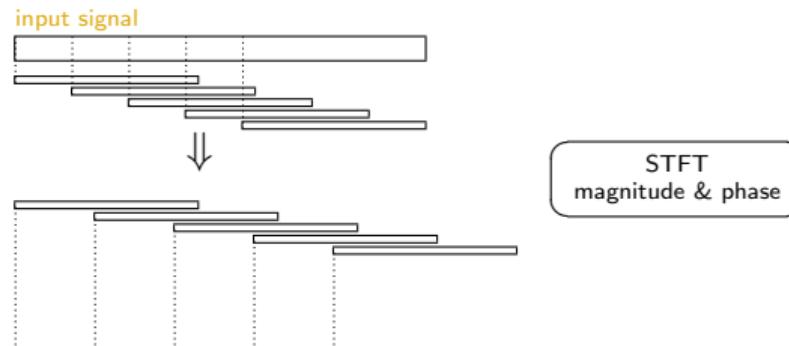
- ① **split input** signal into overlapping blocks
- ② compute **magnitude and phase spectrum** of each block
- ③ change **overlap ratio** between blocks depending on stretch factor
- ④ keep the magnitude, adapt the phase per **bin** to the block's new time stamp



# time stretching

## phase vocoder

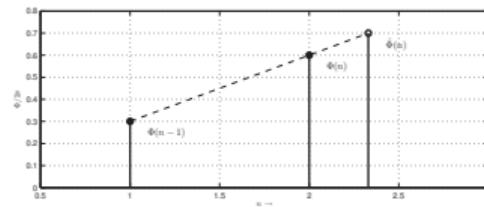
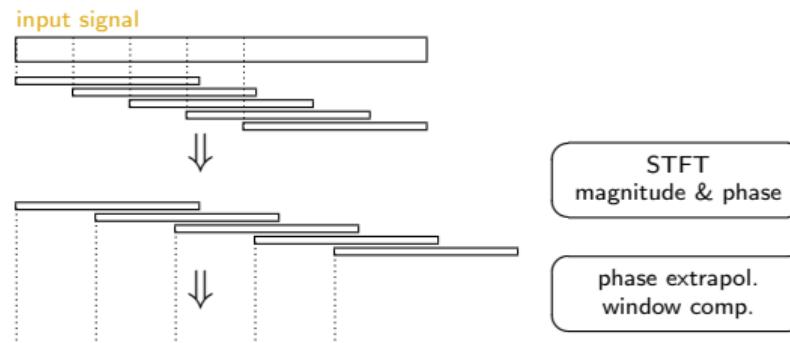
- ① **split input** signal into overlapping blocks
- ② compute **magnitude and phase spectrum** of each block
- ③ **change overlap ratio** between blocks depending on stretch factor
- ④ keep the magnitude, adapt the phase per bin to the block's new time stamp



# time stretching

## phase vocoder

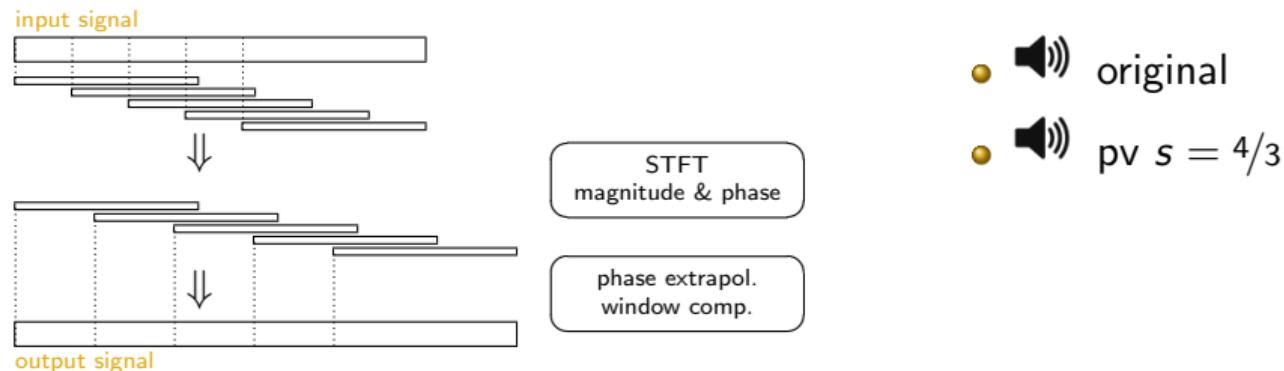
- ① **split input** signal into overlapping blocks
- ② compute **magnitude and phase spectrum** of each block
- ③ **change overlap ratio** between blocks depending on stretch factor
- ④ keep the magnitude, **adapt the phase per bin** to the block's new time stamp



# time stretching

## phase vocoder

- ① **split input** signal into overlapping blocks
- ② compute **magnitude and phase spectrum** of each block
- ③ **change overlap ratio** between blocks depending on stretch factor
- ④ keep the magnitude, **adapt the phase per bin** to the block's new time stamp



# time stretching

frequency reassignment: relation of phase and frequency 1/2

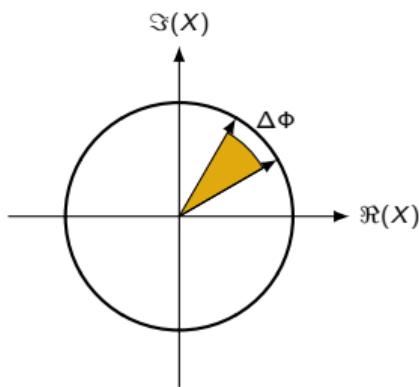
- phasor representation:

- ① sine value is defined by magnitude and phase
- ② decreasing the amplitude  $\Rightarrow$  shorter vector
- ③ increasing the frequency  $\Rightarrow$  increasing speed



# time stretching

frequency reassignment: relation of phase and frequency 2/2



- frequency and phase change closely related:
  - time for full rotation is period length  $T$  with

$$f = \frac{1}{T}$$

- time for fractional rotation  $\Delta\Phi$  is corresponding fraction of period length

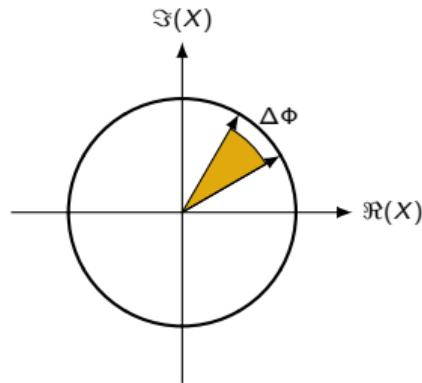
$$f = \frac{\Delta\Phi}{\Delta t}$$

- in other words:

$$\begin{aligned}\Phi(t) &= \omega \cdot t \\ \Rightarrow \frac{d\Phi(t)}{dt} &= \omega = 2\pi f\end{aligned}$$

# time stretching

frequency reassignment: relation of phase and frequency 2/2



- frequency and phase change closely related:
  - time for full rotation is period length  $T$  with

$$f = \frac{1}{T}$$

- time for fractional rotation  $\Delta\Phi$  is corresponding fraction of period length

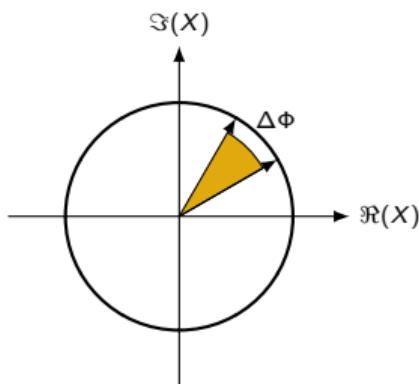
$$f = \frac{\Delta\Phi}{\Delta t}$$

- in other words:

$$\begin{aligned}\Phi(t) &= \omega \cdot t \\ \Rightarrow \frac{d\Phi(t)}{dt} &= \omega = 2\pi f\end{aligned}$$

# time stretching

frequency reassignment: relation of phase and frequency 2/2



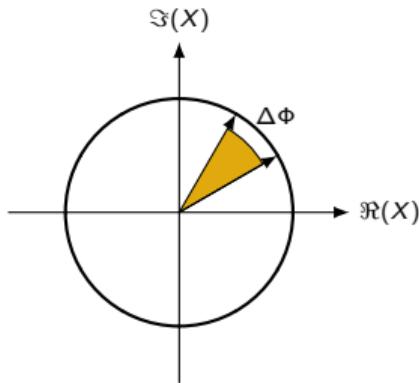
- frequency and phase change closely related:
  - time for full rotation is period length  $T$  with
$$f = \frac{1}{T}$$
  - time for fractional rotation  $\Delta\Phi$  is corresponding fraction of period length
$$f = \frac{\Delta\Phi}{\Delta t}$$

- in other words:

$$\begin{aligned}\Phi(t) &= \omega \cdot t \\ \Rightarrow \frac{d\Phi(t)}{dt} &= \omega = 2\pi f\end{aligned}$$

# time stretching

frequency reassignment: relation of phase and frequency 2/2



- frequency and phase change closely related:
  - time for full rotation is period length  $T$  with
$$f = \frac{1}{T}$$
  - time for fractional rotation  $\Delta\Phi$  is corresponding fraction of period length
$$f = \frac{\Delta\Phi}{\Delta t}$$
  - in other words:

$$\begin{aligned}\Phi(t) &= \omega \cdot t \\ \Rightarrow \frac{d\Phi(t)}{dt} &= \omega = 2\pi f\end{aligned}$$

# time stretching

## frequency reassignment: principles

frequency domain:

- instead of using the bin frequency

$$f(k) = k * \frac{f_S}{\mathcal{K}}$$

- we use the phase of each bin  $\Phi(k, n)$
- to compute the frequency from the phase difference of neighboring blocks

$$\omega_I(k, n) \propto \Phi(k, n) - \Phi(k, n - 1)$$

- $\omega_I(k, n)$  is called **instantaneous frequency** per block per bin

# time stretching

## frequency reassignment: principles

frequency domain:

- instead of using the bin frequency

$$f(k) = k * \frac{f_S}{\mathcal{K}}$$

- we use the phase of each bin  $\Phi(k, n)$
- to compute the frequency from the phase difference of neighboring blocks

$$\omega_I(k, n) \propto \Phi(k, n) - \Phi(k, n - 1)$$

- $\omega_I(k, n)$  is called **instantaneous frequency** per block per bin

# time stretching

frequency reassignment: scaling factor

- instantaneous frequency calculation has to take into account
  - hop size  $\mathcal{H}$
  - sample rate  $f_S$

$$\omega_I(k, n) = \frac{\Delta\Phi_u(k, n)}{\mathcal{H}} \cdot f_S$$

- problem: phase ambiguity

$$\Phi(k, n) = \Phi(k, n) + j \cdot 2\pi$$

⇒ *phase unwrapping*

# time stretching

frequency reassignment: scaling factor

- instantaneous frequency calculation has to take into account
  - hop size  $\mathcal{H}$
  - sample rate  $f_S$

$$\omega_I(k, n) = \frac{\Delta\Phi_u(k, n)}{\mathcal{H}} \cdot f_S$$

- problem: phase ambiguity

$$\Phi(k, n) = \Phi(k, n) + j \cdot 2\pi$$

⇒ *phase unwrapping*

# time stretching

frequency reassignment: scaling factor

- instantaneous frequency calculation has to take into account
  - hop size  $\mathcal{H}$
  - sample rate  $f_S$

$$\omega_I(k, n) = \frac{\Delta\Phi_u(k, n)}{\mathcal{H}} \cdot f_S$$

- problem: phase ambiguity

$$\Phi(k, n) = \Phi(k, n) + j \cdot 2\pi$$

⇒ *phase unwrapping*

# time stretching

frequency reassignment: phase unwrapping

## ① compute unwrapped phase $\Phi_u(k, n)$

- estimate unwrapped bin phase

$$\hat{\Phi}(k, n) = \Phi(k, n - 1) + \underbrace{2\pi k \cdot \frac{\mathcal{H}}{\mathcal{K}}}_{= \omega_k \cdot \frac{\mathcal{H}}{f_s}}$$

- unwrap phase by shifting current phase to estimate's range

$$\Phi_u(k, n) = \hat{\Phi}(k, n) + \text{princarg} [\Phi(k, n) - \hat{\Phi}(k, n)]$$

## ② compute unwrapped phase difference

$$\begin{aligned}\Delta\Phi_u(k, n) &= \Phi_u(k, n) - \Phi(k, n - 1) \\ &= \hat{\Phi}(k, n) + \text{princarg} [\Phi(k, n) - \hat{\Phi}(k, n)] - \Phi(k, n - 1) \\ &= \frac{2\pi k}{\mathcal{K}} \mathcal{H} + \text{princarg} \left[ \Phi(k, n) - \Phi(k, n - 1) - \frac{2\pi k}{\mathcal{K}} \mathcal{H} \right]\end{aligned}$$

# time stretching

frequency reassignment: phase unwrapping

- ① compute unwrapped phase  $\Phi_u(k, n)$

- estimate unwrapped bin phase

$$\hat{\Phi}(k, n) = \Phi(k, n - 1) + \underbrace{2\pi k \cdot \frac{\mathcal{H}}{\mathcal{K}}}_{= \omega_k \cdot \frac{\mathcal{H}}{f_s}}$$

- unwrap phase by shifting current phase to estimate's range

$$\Phi_u(k, n) = \hat{\Phi}(k, n) + \text{princarg} [\Phi(k, n) - \hat{\Phi}(k, n)]$$

- ② compute unwrapped phase difference

$$\begin{aligned}\Delta\Phi_u(k, n) &= \Phi_u(k, n) - \Phi(k, n - 1) \\ &= \hat{\Phi}(k, n) + \text{princarg} [\Phi(k, n) - \hat{\Phi}(k, n)] - \Phi(k, n - 1) \\ &= \frac{2\pi k}{\mathcal{K}} \mathcal{H} + \text{princarg} [\Phi(k, n) - \Phi(k, n - 1) - \frac{2\pi k}{\mathcal{K}} \mathcal{H}]\end{aligned}$$

# time stretching

frequency reassignment: phase unwrapping

- ① compute unwrapped phase  $\Phi_u(k, n)$

- estimate unwrapped bin phase

$$\hat{\Phi}(k, n) = \Phi(k, n - 1) + \underbrace{2\pi k \cdot \frac{\mathcal{H}}{\mathcal{K}}}_{= \omega_k \cdot \frac{\mathcal{H}}{f_s}}$$

- unwrap phase by shifting current phase to estimate's range

$$\Phi_u(k, n) = \hat{\Phi}(k, n) + \text{princarg} [\Phi(k, n) - \hat{\Phi}(k, n)]$$

- ② compute unwrapped phase difference

$$\begin{aligned}\Delta\Phi_u(k, n) &= \Phi_u(k, n) - \Phi(k, n - 1) \\ &= \hat{\Phi}(k, n) + \text{princarg} [\Phi(k, n) - \hat{\Phi}(k, n)] - \Phi(k, n - 1) \\ &= \frac{2\pi k}{\mathcal{K}} \mathcal{H} + \text{princarg} \left[ \Phi(k, n) - \Phi(k, n - 1) - \frac{2\pi k}{\mathcal{K}} \mathcal{H} \right]\end{aligned}$$

# time stretching

frequency reassignment: problems

## • overlapping spectral components

- sinusoidal components often overlap (spectral leakage, several instruments playing the same pitch, ...)
  - ⇒ incorrect phase estimate
  - spectrum should be as sparse as possible, increase STFT length

## • inaccurate phase unwrapping

- unwrapping algorithm is based on assumption of similarity between predicted and measured phase
  - decrease hop size

# time stretching

frequency reassignment: problems

## • overlapping spectral components

- sinusoidal components often overlap (spectral leakage, several instruments playing the same pitch, ...)   
⇒ incorrect phase estimate
- spectrum should be as sparse as possible, increase STFT length

## • inaccurate phase unwrapping

- unwrapping algorithm is based on assumption of similarity between predicted and measured phase
- decrease hop size

# time stretching

frequency reassignment: problems

## • overlapping spectral components

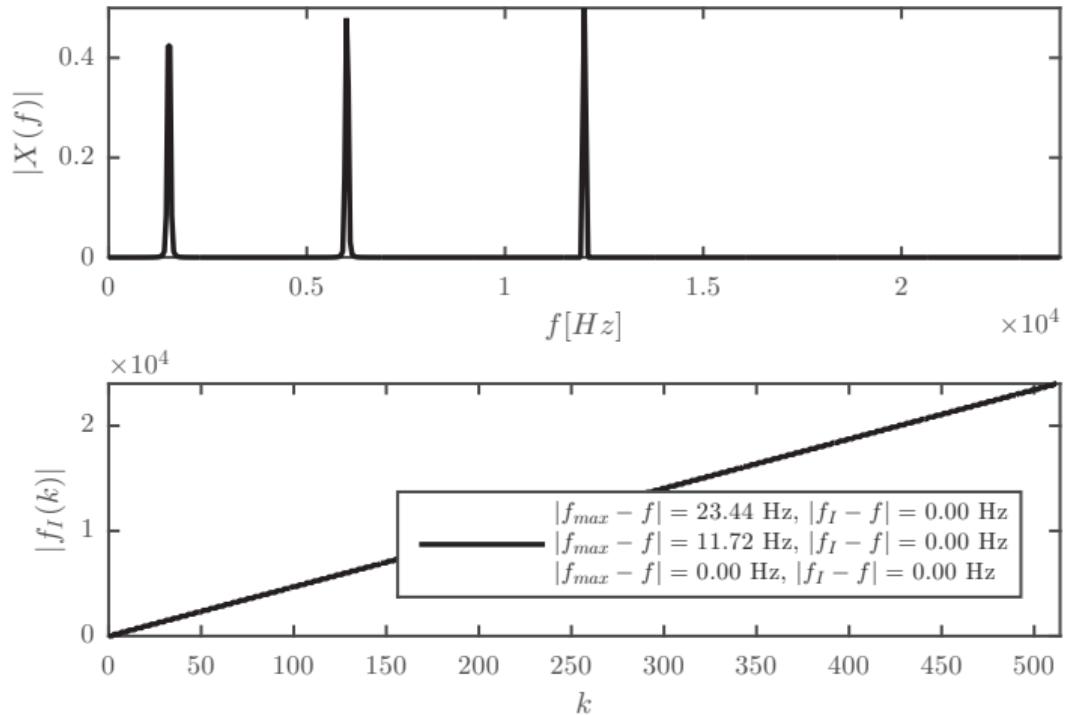
- sinusoidal components often overlap (spectral leakage, several instruments playing the same pitch, ...)   
⇒ incorrect phase estimate
  - spectrum should be as sparse as possible, increase STFT length

## • inaccurate phase unwrapping

- unwrapping algorithm is based on assumption of similarity between predicted and measured phase
  - decrease hop size

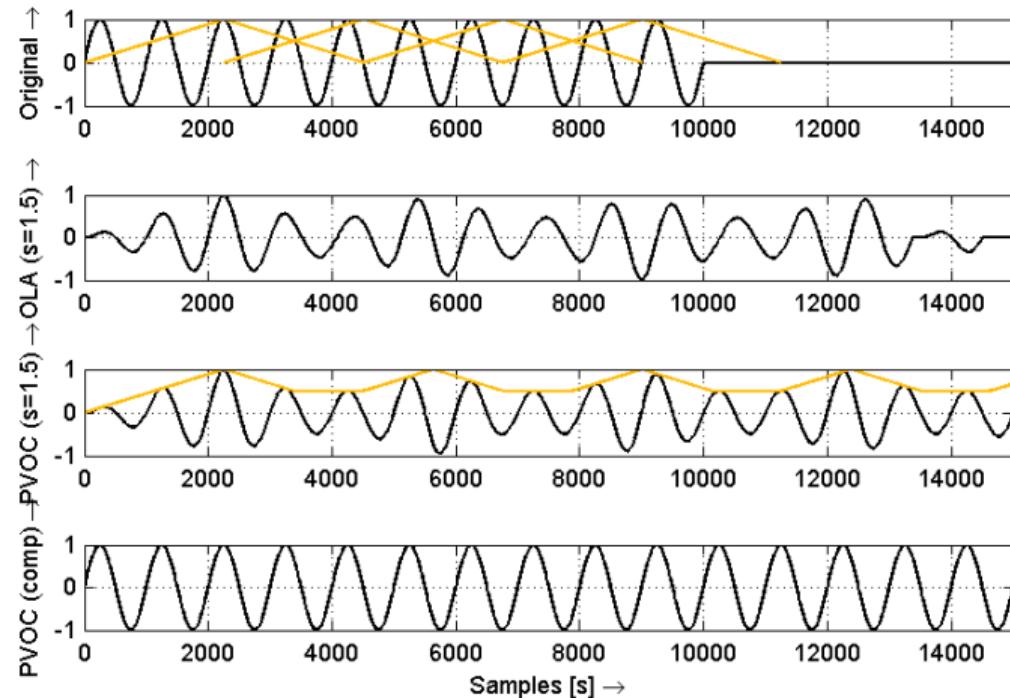
# time stretching

frequency reassignment: example



# time stretching

## phase vocoder window compensation



# time stretching

## phase vocoder — properties & artifacts

- **advantages:**

- allows *Polyphonic input* (assumption: no overlapping harmonics)
- absolute *timing stability* (i.e., sample resolution)

- **disadvantages:**

- *low granularity* — FFT block size
- artifacts:
  - ① phasing
  - ② transient smearing/doubling

# time stretching

## phase vocoder — properties & artifacts

- **advantages:**

- allows *Polyphonic input* (assumption: no overlapping harmonics)
- absolute *timing stability* (i.e., sample resolution)

- **disadvantages:**

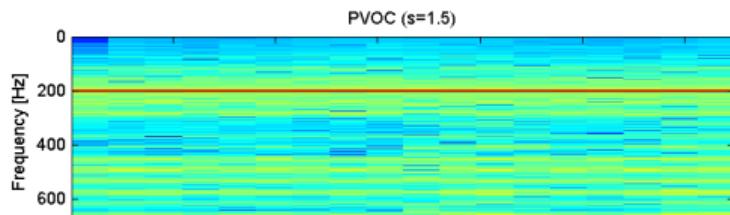
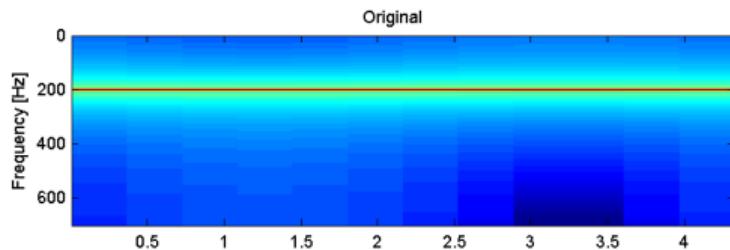
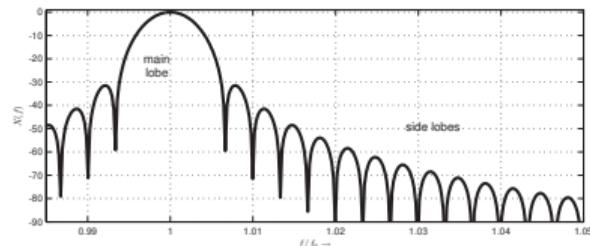
- *low granularity* — FFT block size

- artifacts:

- ① phasing
  - ② transient smearing/doubling

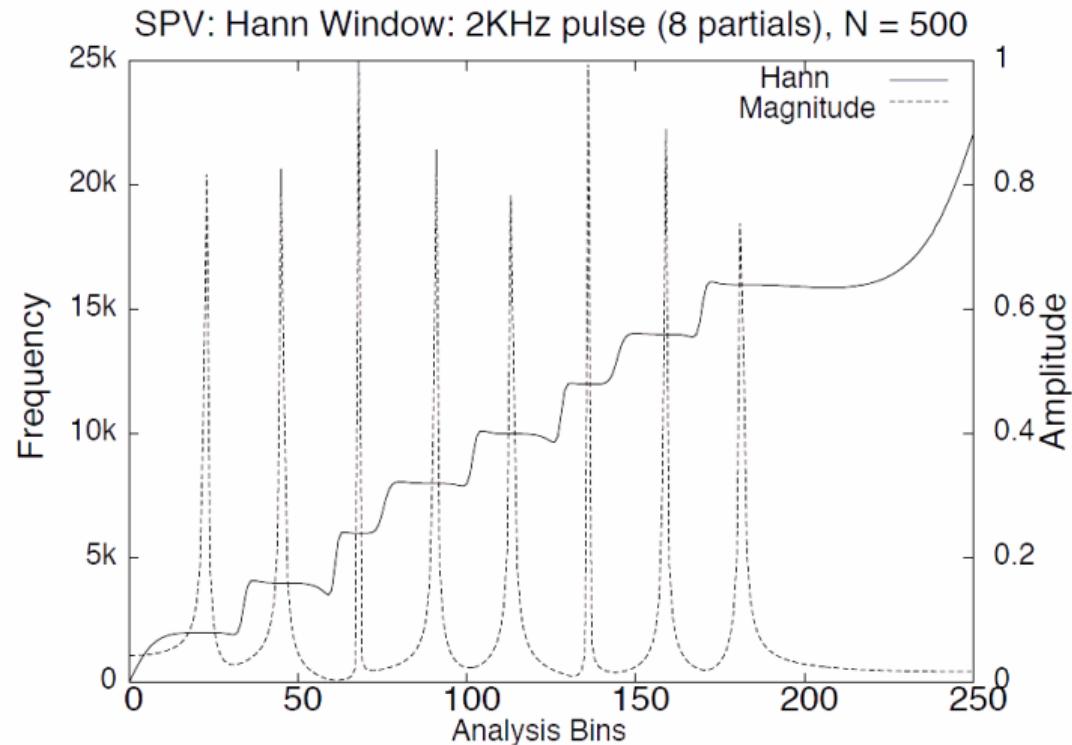
# time stretching

phase vocoder artifacts: phasing — spectral leakage 1/3



# time stretching

phase vocoder artifacts: phasing — spectral leakage 2/3



# time stretching

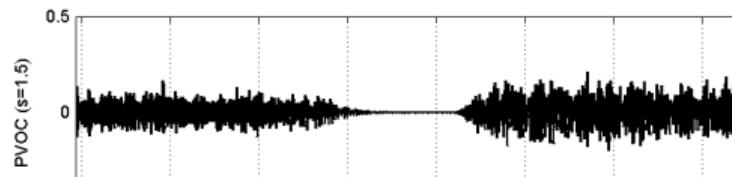
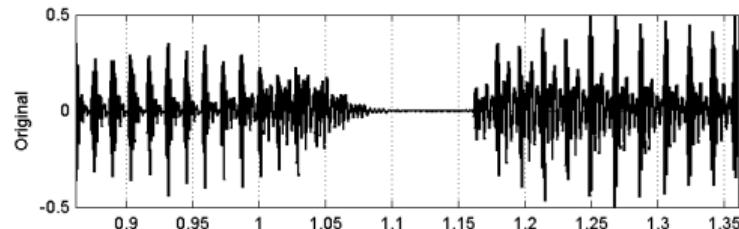
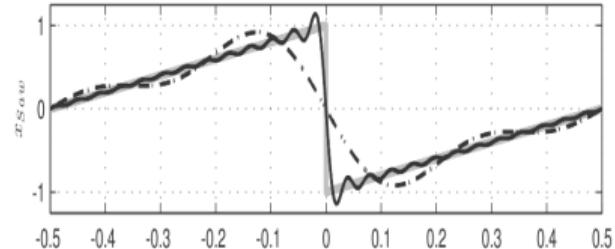
phase vocoder artifacts: phasing — spectral leakage 3/3

⇒ use *frequency reassignment* for grouping and phase sync

- original
- $pv\ s = 4/3$
- grouped phase

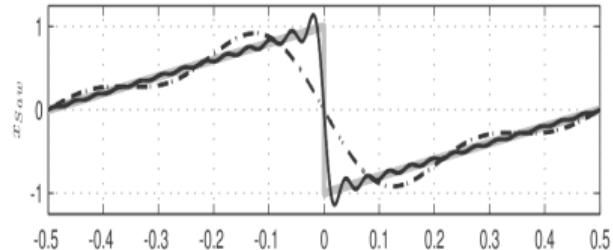
# time stretching

phase vocoder artifacts: phasing — unsynced harmonics



# time stretching

phase vocoder artifacts: phasing — unsynced harmonics



⇒ use *harmonic analysis* for grouping and phase sync

- original
- pv  $s = 4/3$
- synced phase

# time stretching

phase vocoder artifacts: interchannel phasing

- phase estimation between channels slightly off due to
  - numerical inaccuracies (cumulative!)
  - overlapping frequency components

⇒ change in spatial image

-  original
-  pv  $s = 3/2$

# time stretching

phase vocoder artifacts: interchannel phasing

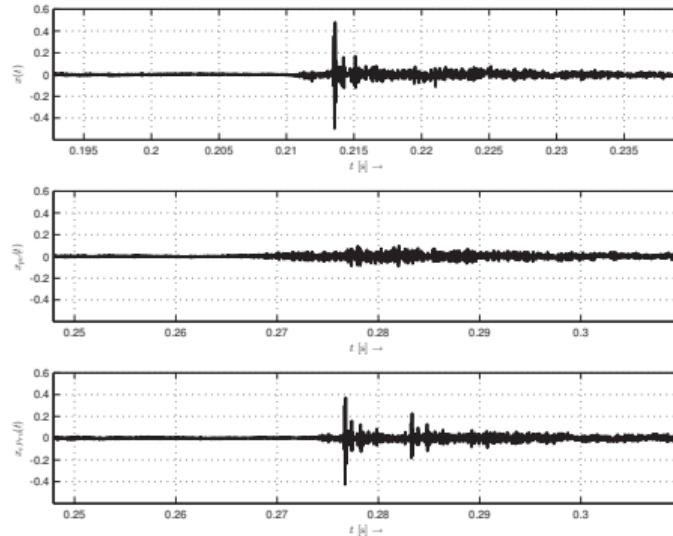
- phase estimation between channels slightly off due to
  - numerical inaccuracies (cumulative!)
  - overlapping frequency components

⇒ change in spatial image

-  original
-  pv  $s = 3/2$

# time stretching

phase vocoder artifacts: transient smearing



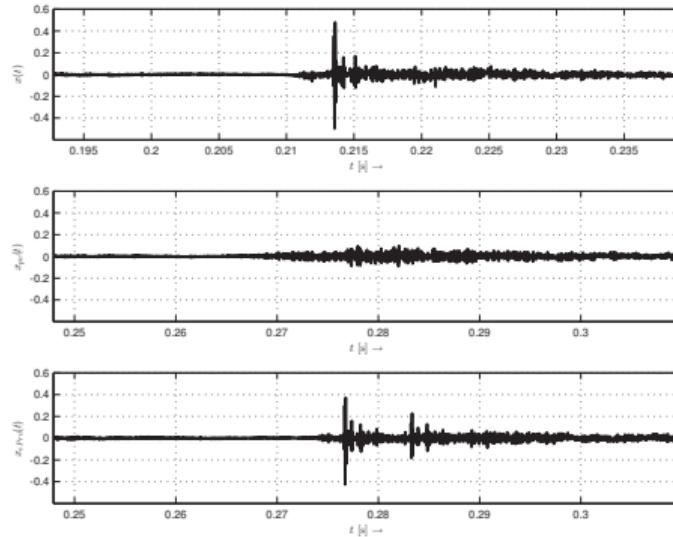
- original
- pv  $s = 3/2$

⇒ detect transients and *reset phase per bin*

- original
- pv  $s = 4/3$
- phase reset

# time stretching

phase vocoder artifacts: transient smearing



- original
- pv  $s = 3/2$

⇒ detect transients and *reset phase per bin*

- original
- pv  $s = 4/3$
- phase reset

# time stretching

## inherent problems

- stretching the audio data can lead to “**non-natural**” results
- examples
  - tempo dependent *timing variations*
  - other performance related aspects may get inappropriate lengths and speed: *vibrato, tremolo, glissando*

# time stretching

## inherent problems

- stretching the audio data can lead to “**non-natural**” results
- **examples**
  - tempo dependent *timing variations*
  - other performance related aspects may get inappropriate lengths and speed: *vibrato, tremolo, glissando*

# time stretching

## inherent problems

- stretching the audio data can lead to “**non-natural**” results
- **examples**
  - tempo dependent *timing variations*
  - other performance related aspects may get inappropriate lengths and speed: *vibrato, tremolo, glissando*

# pitch shifting

## standard approach

- **definition:**

change pitch without changing tempo

- **method:**

combine stretching and *sample rate conversion* (interpolation)

- ① change length with stretching
- ② resample to compensate for length difference

- **implementation:** differentiate "external" and "internal" parameters

- *extern*: stretch  $s_e$  and pitch  $p_e$
- *intern*: stretch  $s_i$  and resample  $r_i$

- **example:** pitch shift factor  $p = 4/3$

- ③ *time stretch* (increase length/decrease tempo)  $s = 4/3$
  - ④ *resample* (decrease length, increase pitch)  $r = 3/4$
- ⇒ audio:



# pitch shifting

## standard approach

- **definition:**

change pitch without changing tempo

- **method:**

combine stretching and *sample rate conversion* (interpolation)

- ① change length with stretching
- ② resample to compensate for length difference

- **implementation:** differentiate “external” and “internal” parameters

- *extern*: stretch  $s_e$  and pitch  $p_e$
- *intern*: stretch  $s_i$  and resample  $r_i$

- **example:** pitch shift factor  $p = 4/3$

- ① time stretch (increase length/decrease tempo)  $s = 4/3$
  - ② resample (decrease length, increase pitch)  $r = 3/4$
- audio:



• OLA



• phase vocoder

# pitch shifting

## standard approach

- **definition:**

change pitch without changing tempo

- **method:**

combine stretching and *sample rate conversion* (interpolation)

- 1 change length with stretching
- 2 resample to compensate for length difference

- **implementation:** differentiate “external” and “internal” parameters

- *extern*: stretch  $s_e$  and pitch  $p_e$
- *intern*: stretch  $s_i$  and resample  $r$ ;

- **example:** pitch shift factor  $p = 4/3$

- 1 *time stretch* (increase length/decrease tempo)  $s = 4/3$
- 2 *resample* (decrease length, increase pitch)  $r = 3/4$

⇒ audio:



# pitch shifting

## standard approach

- **definition:**

change pitch without changing tempo

- **method:**

combine stretching and *sample rate conversion* (interpolation)

- ① change length with stretching
- ② resample to compensate for length difference

- **implementation:** differentiate “external” and “internal” parameters

- *extern*: stretch  $s_e$  and pitch  $p_e$
- *intern*: stretch  $s_i$  and resample  $r$ ;

- **example:** pitch shift factor  $p = 4/3$

- ① *time stretch* (increase length/decrease tempo)  $s = 4/3$
  - ② *resample* (decrease length, increase pitch)  $r = 3/4$
- ⇒ audio:



# pitch shifting

standard approach: stretch, pitch, resample factor examples

- extern:  $s_e = 1$   $p_e = 2$

# pitch shifting

standard approach: stretch, pitch, resample factor examples

- extern:  $s_e = 1$   $p_e = 2$   
intern:  $s_i = 2$   $r_i = \frac{1}{2}$

# pitch shifting

standard approach: stretch, pitch, resample factor examples

- extern:  $s_e = 1$   $p_e = 2$   
intern:  $s_i = 2$   $r_i = \frac{1}{2}$
- extern:  $s_e = 1$   $p_e = \frac{4}{3}$

# pitch shifting

standard approach: stretch, pitch, resample factor examples

- extern:  $s_e = 1$   $p_e = 2$

intern:  $s_i = 2$   $r_i = \frac{1}{2}$

- extern:  $s_e = 1$   $p_e = \frac{4}{3}$

intern:  $s_i = \frac{4}{3}$   $r_i = \frac{3}{4}$

# pitch shifting

standard approach: stretch, pitch, resample factor examples

- extern:  $s_e = 1$   $p_e = 2$   
intern:  $s_i = 2$   $r_i = \frac{1}{2}$
- extern:  $s_e = 1$   $p_e = \frac{4}{3}$   
intern:  $s_i = \frac{4}{3}$   $r_i = \frac{3}{4}$
- extern:  $s_e = \frac{1}{2}$   $p_e = 2$

# pitch shifting

standard approach: stretch, pitch, resample factor examples

- extern:  $s_e = 1$   $p_e = 2$   
intern:  $s_i = 2$   $r_i = \frac{1}{2}$
- extern:  $s_e = 1$   $p_e = \frac{4}{3}$   
intern:  $s_i = \frac{4}{3}$   $r_i = \frac{3}{4}$
- extern:  $s_e = \frac{1}{2}$   $p_e = 2$   
intern:  $s_i = 1$   $r_i = \frac{1}{2}$

# pitch shifting

standard approach: stretch, pitch, resample factor examples

- extern:  $s_e = 1$   $p_e = 2$   
intern:  $s_i = 2$   $r_i = \frac{1}{2}$
- extern:  $s_e = 1$   $p_e = \frac{4}{3}$   
intern:  $s_i = \frac{4}{3}$   $r_i = \frac{3}{4}$
- extern:  $s_e = \frac{1}{2}$   $p_e = 2$   
intern:  $s_i = 1$   $r_i = \frac{1}{2}$
- extern:  $s_e = 2$   $p_e = 2$

# pitch shifting

standard approach: stretch, pitch, resample factor examples

- extern:  $s_e = 1$   $p_e = 2$   
intern:  $s_i = 2$   $r_i = \frac{1}{2}$
- extern:  $s_e = 1$   $p_e = \frac{4}{3}$   
intern:  $s_i = \frac{4}{3}$   $r_i = \frac{3}{4}$
- extern:  $s_e = \frac{1}{2}$   $p_e = 2$   
intern:  $s_i = 1$   $r_i = \frac{1}{2}$
- extern:  $s_e = 2$   $p_e = 2$   
intern:  $s_i = 4$   $r_i = \frac{1}{2}$

# pitch shifting

## frequency domain approach

### ① STFT

- ② magnitude and phase
- ③ magnitude and instantaneous frequency
- ④ resample both magnitude and frequency spectrum acc. to pitch factor
- ⑤ magnitude and phase
- ⑥ complex spectrum
- ⑦ IFFT and OLA

# pitch shifting

## frequency domain approach

- ➊ STFT
- ➋ magnitude and phase
- ➌ magnitude and instantaneous frequency
- ➍ resample both magnitude and frequency spectrum acc. to pitch factor
- ➎ magnitude and phase
- ➏ complex spectrum
- ➐ IFFT and OLA

# pitch shifting

## frequency domain approach

- ① STFT
- ② magnitude and phase
- ③ magnitude and instantaneous frequency
- ④ resample both magnitude and frequency spectrum acc. to pitch factor
- ⑤ magnitude and phase
- ⑥ complex spectrum
- ⑦ IFFT and OLA

# pitch shifting

## frequency domain approach

- ➊ STFT
- ➋ magnitude and phase
- ➌ magnitude and instantaneous frequency
- ➍ resample both magnitude and frequency spectrum acc. to pitch factor
- ➎ magnitude and phase
- ➏ complex spectrum
- ➐ IFFT and OLA

# pitch shifting

## frequency domain approach

- ➊ STFT
- ➋ magnitude and phase
- ➌ magnitude and instantaneous frequency
- ➍ resample both magnitude and frequency spectrum acc. to pitch factor
- ➎ magnitude and phase
- ➏ complex spectrum
- ➐ IFFT and OLA

# pitch shifting

## frequency domain approach

- ➊ STFT
- ➋ magnitude and phase
- ➌ magnitude and instantaneous frequency
- ➍ resample both magnitude and frequency spectrum acc. to pitch factor
- ➎ magnitude and phase
- ➏ complex spectrum
- ➐ IFFT and OLA

# pitch shifting

## frequency domain approach

- ① STFT
- ② magnitude and phase
- ③ magnitude and instantaneous frequency
- ④ resample both magnitude and frequency spectrum acc. to pitch factor
- ⑤ magnitude and phase
- ⑥ complex spectrum
- ⑦ IFFT and OLA

# pitch shifting

formant preservation: time domain

## • idea

- signal is pulse train filtered by transfer function
  - pulse train determines fundamental frequency
  - transfer function determines formant shape/timbre characteristics

## • approach

- change grain/pulse distance
- grain “content” not modified  
→ freq domain not modified

⇒ “same” spectrum,  
different pitch

# pitch shifting

formant preservation: time domain

## • idea

- signal is pulse train filtered by transfer function
  - pulse train determines fundamental frequency
  - transfer function determines formant shape/timbre characteristics

## • approach

- change grain/pulse distance
- grain “content” not modified  
→ freq domain not modified

⇒ “same” spectrum,  
different pitch

# pitch shifting

formant preservation: time domain

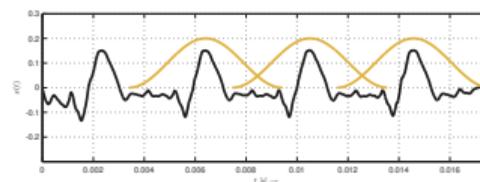
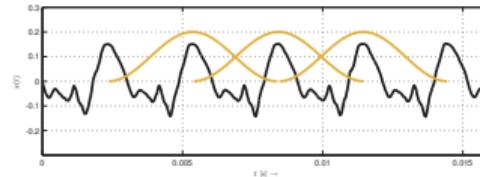
## • idea

- signal is pulse train filtered by transfer function
  - pulse train determines fundamental frequency
  - transfer function determines formant shape/timbre characteristics

## • approach

- change grain/pulse distance
- grain “content” not modified  
→ freq domain not modified

⇒ “same” spectrum,  
different pitch



# pitch shifting

formant preservation: time domain

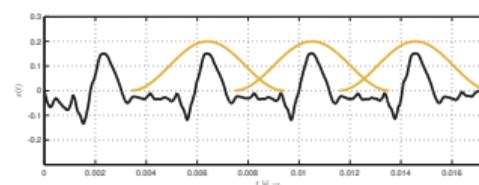
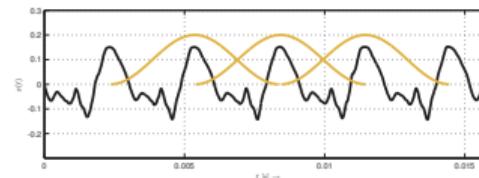
## • idea

- signal is pulse train filtered by transfer function
  - pulse train determines fundamental frequency
  - transfer function determines formant shape/timbre characteristics

## • approach

- change grain/pulse distance
- grain “content” not modified  
→ freq domain not modified

⇒ “same” spectrum,  
different pitch



# pitch shifting

formant preservation: frequency domain

- **idea**

- preserve spectral envelope

- **approach**

- 1 measure spectral envelope
  - 2 apply inverse envelope (whitening)
  - 3 pitch shift
  - 4 apply spectral envelope

# pitch shifting

formant preservation: frequency domain

- **idea**

- preserve spectral envelope

- **approach**

- ① measure spectral envelope
  - ② apply inverse envelope (whitening)
  - ③ pitch shift
  - ④ apply spectral envelope

# pitch shifting

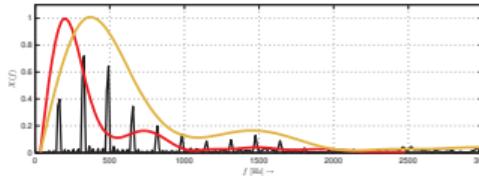
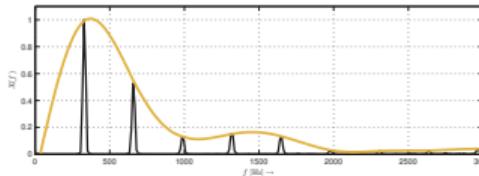
formant preservation: frequency domain

- **idea**

- preserve spectral envelope

- **approach**

- 1 measure spectral envelope
- 2 apply inverse envelope (whitening)
- 3 pitch shift
- 4 apply spectral envelope



# pitch shifting

## spectral envelope estimate

### ● approaches

- LPC coefficients
- spectral maxima

### ● potential issues

- *Polyphonic input* audio:  
'superposition' of envelopes
- *Very high/low pitch factors*: high frequency boost/cut
- *estimate resolution*
  - too coarse → loss of timbre characteristics
  - too fine → impress pitch characteristics (harmonic pattern) on spectrum

# pitch shifting

## spectral envelope estimate

- **approaches**

- LPC coefficients
- spectral maxima

- **potential issues**

- *polyphonic input* audio:  
'superposition' of envelopes
- *very high/low pitch factors*: high frequency boost/cut
- *estimate resolution*
  - too coarse → loss of timbre characteristics
  - too fine → impress pitch characteristics (harmonic pattern) on spectrum

# pitch shifting

## audio examples

	OLA	PSOLA	PVOC
orig			
resample $p = 4/3$			
formant $p = 4/3$			

# summary

## time stretching and pitch shifting

- pitch stretching and pitch shifting are largely equivalent algorithms
  - same artifacts
  - same workload
- monophonic time-stretching with PSOLA-based approaches is
  - easier to solve
  - has very bad artifacts if pitch tracker is off
- polyphonic time-stretching with PV-based approaches is
  - complicated due to tradeoffs (e.g., frequency vs time resolution)
- general challenges:
  - noisy and transient signals
  - resulting timbre changes
  - perceived naturalness of result
  - time resolution/accuracy due to blocked processing